

---

---

# DS595/CS525 FALL

## Project 3 - Deep Q-learning

— 10/3/2019 —

---

---

# Outline

- Introduction
  - Game Playing : Breakout
- Deep Reinforcement Learning
  - Deep Q-Learning (DQN)
  - Improvements to DQN
- Grading & Format
  - Grading Policy
  - Code Format
  - Submission
- Google Cloud Platform & Pytorch Tutorial *(12-13 hours on GPU)*

# Introduction

## Environment


### Breakout



- Get average reward  $\geq 40$  in 100 episodes
- With **OpenAI's Atari wrapper** & reward clipping
  - We will unclip the reward when testing

# Introduction

## Training Plot



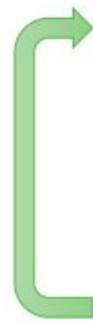
- X-axis : number of training steps
- Y-axis : average clipped reward in last 30 episodes

# Deep Reinforcement Learning

## Deep Q-Learning (DQN)

“classic” deep Q-learning algorithm:


Replay buffer

- 
1. take some action  $\mathbf{a}_i$  and observe  $(\mathbf{s}_i, \mathbf{a}_i, \mathbf{s}'_i, r_i)$ , add it to  $\mathcal{B}$
  2. sample mini-batch  $\{\mathbf{s}_j, \mathbf{a}_j, \mathbf{s}'_j, r_j\}$  from  $\mathcal{B}$  uniformly
  3. compute  $y_j = r_j + \gamma \max_{\mathbf{a}'_j} Q_{\phi'}(\mathbf{s}'_j, \mathbf{a}'_j)$  using *target* network  $Q_{\phi'}$
  4.  $\phi \leftarrow \phi - \alpha \sum_j \frac{dQ_{\phi}}{d\phi}(\mathbf{s}_j, \mathbf{a}_j)(Q_{\phi}(\mathbf{s}_j, \mathbf{a}_j) - y_j)$
  5. update  $\phi'$ : copy  $\phi$  every  $N$  steps

# Deep Reinforcement Learning

## Deep Q-Learning (DQN)

- The action should act  $\epsilon$ -greedily
  - Random action with probability  $\epsilon$
- Linearly decline  $\epsilon$  from 1.0 to some small value, say 0.025
  - Decline per step
- Hyperparameters (just suggestion)
  - Replay Buffer Memory Size 10000 (*deque*)
  - Start to learn 5000
  - Perform Update Target Network Step 5000
  - Learning Rate 1.5e-4, Batch Size 32
  - Adam
  - Huber Loss (*F.smooth\_l1\_loss*)
  - Clip gradients between (-1,1)



Green is the Huber loss and blue is the quadratic loss (Wikipedia)

$$L_{\delta}(a) = \begin{cases} \frac{1}{2}a^2 & \text{for } |a| \leq \delta, \\ \delta(|a| - \frac{1}{2}\delta), & \text{otherwise.} \end{cases}$$

# Deep Reinforcement Learning

## Why Reward is clipped

- Performing the same action for 4 frames
  - To use data more efficiently

```
env = Environment('BreakoutNoFrameskip-v4', '', atari_wrapper=True, test=True)
```

- Reward may be up to 4
  - If positive, clip to 1 → reduce variance

```
env.step(0)[0].shape  
(84, 84, 4)
```

- How to see your unclipped reward

```
env.step(0)[0].dtype  
dtype('uint8')
```

1. Use the *test* function
2. Turn off the *clip\_reward* option of your environment and do the clipping by yourself.

# Deep Reinforcement Learning

## Improvements to DQN

- Double Q-Learning
- Dueling Network
- Prioritized Replay Memory
- Noisy DQN
- Distributional DQN

<https://arxiv.org/pdf/1710.02298.pdf>



# Deep Reinforcement Learning

## Other Training Tips

- [How to use Pytorch](#)
- [Official DQN Pytorch Tutorial](#)
- [DQN Tutorial on Medium](#)
- [Official DQN paper](#)

# Grading & Format

## Grading Policy

- Python code (20 points)
- Trained Model (50 points)
  - Get averaging reward  $\geq 40$  in 100 episodes in **Breakout**
  - With **OpenAI's Atari wrapper** & reward clipping
    - We will unclip the reward when testing
- PDF Report (30 points)
  - Describe your DQN model
  - Plot the learning curve (*training steps can be defined by yourself*)
    - X-axis: number of training steps
    - Y-axis: average clipped reward in last 30 episodes

# Grading & Format

## Code Format

- Please download the sample files from [github](#)
- Follow the instructions in README to install required packages
- **Six** functions you should implement in [agent\\_dqn.py](#)
  1. `__init__(self, env, args)`
  2. `init_game_setting(self)`
  3. `make_action(self, state, test)`
  4. `train(self)`
  5. `push(self)`
  6. `repaly_buffer(self)`
- **DO NOT** add any parameter in `__init__()`, `init_game_setting()` and `make_action()`
- You can change the seed
- You can add new functions in the [agent\\_dqn.py](#)

# Grading & Format

## Code Format

- **Two** functions you should implement in `dqn_model.py`
  1. `__init__(self)`
  2. `forward(self, x)`
- You can add parameters in these two functions
- You can add new functions in the `dqn_model.py`
- You can add your arguments in `argument.py`
- **Please don't revise `test.py`, `environment.py` and `agent.py`**

# Grading & Format

## Deliverables

- Deadline: **Thursday 17/10/2019 23:59**
- Your submission **MUST** have following files
  - `agent_dqn.py`, `dqn_model.py`, `argument.py`, `atari_wrapper.py`
  - `[saved_model_file]`
  - `report.pdf`
  - README (optional)
  - `download.sh` (optional)
  - other files you need
- If your model is too large for canvas, upload it to a cloud space and write `download.sh` to download the model

# Grading & Format

## Package

- Please use Python3
- The TA will execute 'python test.py --test\_dqn' to run your code on **ubuntu+GPU**
- The execution for the model should be done within 20 minutes, excluding model download
- Allowed packages
  - a. PyTorch
  - b. Numpy
  - c. Scipy
  - d. Pandas
  - e. Python Standard Lib
- If you use other packages, please ask for permission first

# Google Cloud Platform

- [How to use Google Cloud Platform](#)
- [How to use Pytorch on GPU](#)
- [Naive Pytorch tutorial](#)

# Related Materials

- Course & Tutorial:
  - [Berkeley Deep Reinforcement Learning, Fall 2017](#)
  - [David Silver RL course](#)
  - [Nips 2016 RL tutorial](#)
- Blog:
  - [Andrej Karpathy's blog](#)
  - [Arthur Juliani's Blog](#)
- Text Book:
  - [Reinforcement Learning: An Introduction](#)
- Repo:
  - <https://github.com/williamFalcon/DeepRLHacks>